

Part V

Memory and pointers

Outline

19 Pointers

20 Memory

21 Function wrappers

Pointers

- A variable `x` stored in memory is stored at a specific place in memory (its address)
- If we store that address in a variable `p`, then `p` is a pointer to `x`

```
int x = 5;  
int* p = &x; // Store address of x in p  
std::cout << *p << std::endl; // Prints value of x  
*p = 6; // Changes x  
std::cout << x; // Prints 6
```

- Note the address-of operator `&`.
- This is different from the pass-by-reference modifier `&` seen earlier.

Declaring multiple pointers

- When declaring multiple variables:

```
int* a,b,c;
```

creates an integer pointer **a** and *integers* **b** and **c**

- A clearer way to write the same thing is

```
int *a, b, c;
```

or

```
int *a;  
int b,c;
```

- It may help to think of the second example as declaring: “A variable which, when dereferenced, gives an integer.”
- Alternatively, just avoid this confusion altogether by using the second or third ways of writing this.

Aliasing

- Through the use of pointers and references, a single variable in memory can be changed using multiple labels
- This can cause programmer errors, if variables are changed in non-obvious ways: just because there is no `x=` doesn't mean that `x` is constant:

```
int f(int& x, int &y) {  
    const int a = x;  
    y = y+1;  
    const int b = x;  
}  
f(s, s); // Would lead to a != b within f
```

- Note that not even using `const` on `x` enforces this

Const-pointers

- Const-ness extends naturally to pointers
- However, some clarification may be needed:

```
const int x = 0;
int y;
const int z = 1;
const int* p; // Pointer to a constant integer
*p = 5; // Compiler-error
p = &x; // OK
p = &y; // OK - but y cannot be changed through p
int* const q = &y; // Constant-pointer to a non-constant
                  integer
*q = 4; // OK - changes y
q = &z; // Compiler error - would change q

const int* const r = &x;
*r = 3; // Compile-error
r = &z; // Compile-error

// Multiple constant ptrs to const vars
const int *const s = &x, *const t = &y;
```

Outline

19 Pointers

20 Memory


21 Function wrappers

Heap and Stack

- The stack is a Last-In-First-Out structure (imagine a stack of plates)
- The stack contains the local data used by all functions currently executing
- When calling a function, its parameters are copied onto the stack, As local variables are allocated, they are also placed on the stack
- When exiting a function:
 - its return value is copied out
 - all local variables are deleted
 - and the function is taken off the stack
- This leaves the calling function at the top, so it can continue with evaluation
- The stack has limited size, so very large arrays can't be stored on it

Stack simulation

```
1  int f(int y){  
2      int z = 3*y*y + 4*y;  
3      return z;  
4  }  
5  
6  int main(void){  
7      int x = f(5);  
8  }
```




Line Counter:

6

Stack simulation

```
1  int f(int y){  
2      int z = 3*y*y + 4*y;  
3      return z;  
4  }  
5  
6  int main(void){  
7      int x = f(5);  
8  }
```




Line Counter:

7

Stack simulation

```
1  int f(int y){  
2    int z = 3*y*y + 4*y;  
3    return z;  
4  }  
5  
6  int main(void){  
7    int x = f(5);  
8  }
```




Line Counter:

1

Stack simulation

```
1  int f(int y){  
2      int z = 3*y*y + 4*y;  
3      return z;  
4  }  
5  
6  int main(void){  
7      int x = f(5);  
8  }
```




Line Counter:

1

Stack simulation

```
1  int f(int y){  
2      int z = 3*y*y + 4*y;  
3      return z;  
4  }  
5  
6  int main(void){  
7      int x = f(5);  
8  }
```




Line Counter:

2

Stack simulation

```
1  int f(int y){  
2      int z = 3*y*y + 4*y;  
3      return z;  
4  }  
5  
6  int main(void){  
7      int x = f(5);  
8  }
```




Line Counter:

2

Stack simulation


```
1  int f(int y){  
2    int z = 3*y*y + 4*y;  
3    return z;  
4  }  
5  
6  int main(void){  
7    int x = f(5);  
8  }
```



Line Counter:
3

Stack simulation

```
1  int f(int y){  
2    int z = 3*y*y + 4*y;  
3    return z;  
4  }  
5  
6  int main(void){  
7    int x = f(5);  
8  }
```




Line Counter:

4

Stack simulation

```
1 int f(int y){  
2   int z = 3*y*y + 4*y;  
3   return z;  
4 }  
5  
6 int main(void){  
7   int x = f(5);  
8 }
```



Line Counter:

7

Stack simulation

```
1  int f(int y){  
2    int z = 3*y*y + 4*y;  
3    return z;  
4  }  
5  
6  int main(void){  
7    int x = f(5);  
8  }
```

int x=95
&main

Line Counter:

7

Heap

- The heap is (essentially) the rest of your computer's memory
- Can store objects of arbitrary size
- Allocation/deletion of objects can be done in any order
- Allocation has some time/memory overhead, so usually good to allocate a few large blocks of memory, rather than many small blocks.
- Objects are not automatically deleted once they go out of scope.
- This is useful because objects persist without being explicitly copied, saving the time taken to copy data.
- This is a problem because you have to remember to release the memory explicitly, otherwise your program's memory usage may grow too large.

Memory allocation (heap)

- We can request an amount of memory from the OS and have it return a pointer to a contiguous (uninitialized) block of memory.

```
int* p = new int; // Allocate memory for a single integer
*p = 8; // Dereferencing
std::cout << *p << std::endl;
delete p; // Free memory
```

- The memory remains allocated until explicitly **deleted** or the program exits
- (It is up to the OS to free on program exit; Linux and Windows do, some embedded OSs may not.)
- So, it is possible to leak memory by not freeing a pointer to it when you have finished with the memory.
- Copying pointers does not copy the associated memory.

Array allocation (heap)

- We can allocate memory for a block of variable size

```
int s = 100;  
int* p = new int[s];  
for(int i=0 ; i < s ; i++){  
    p[i] = i;  
}  
delete[] p; // Free all 100 ints
```

Heap memory available outside fn

- The following is valid because memory allocated on the heap is not freed automatically:

```
int* allocateIntMem(const int a){  
    int* ptr = new int[a];  
    return ptr;  
}  
  
int* myData = allocateIntMem(10);  
myData[0] = 6;  
myData[9] = 8;  
delete[] myData;
```

- `delete` frees the memory pointed to by `myData`, but `myData` still holds its original value.
- Any attempt to access memory through `myData` will be invalid hereafter, and lead to a seg-fault.

nullptr

- There is a C++-literal: `nullptr` usually used for an undefined pointer
- Trying to dereference this will lead to a segmentation-fault
- Usefully, `delete nullptr;` is valid, and does nothing
- This avoids the need for constructs such as

```
if( p != nullptr ){
    delete p;
}
```

or equivalently

```
if( p ){
    delete p;
}
```

- Note also that short-cut evaluation means that expressions such as

```
if( p && *p == 4 )
```

are well-defined since `p` will not be dereferenced if it is `nullptr`

Pointer arithmetic

- Note that the square bracket notation is equivalent to dereferencing:
`p[0]` and `*p` refer to the same thing.
- Arithmetic operations on a pointer are equivalent to `[]`:
`p[1]` and `*(p+1)` are equivalent.
- Modifying a pointer is acceptable:

```
for(int i=0 ; i < 10 ; i++){  
    *p = i;  
    p = p+1;  
}
```

- Adding 1 to a pointer advances it to point to the next item in memory of the same type.
- Note that this causes problems if you later try to **delete** a modified pointer. **delete** must be called on the original pointer returned by **new**.

Aliasing

- It is very easy to end up with aliasing effects again:

```
int* p = new int[N];  
int* q = p;  
q[10] = 10;  
std::cout << p[10] << std::endl; // prints 10
```

- And multiple deletes on the same pointer are invalid:

```
delete[] p;  
q[5] = 10; // Invalid  
delete[] q; // Invalid
```

Smart pointers

- To ensure that allocated memory is always freed, consider a `shared_ptr`.
- This obeys scope, but can be copied, and the object pointed to is destroyed when the last pointer to it is deleted.

```
std::shared_ptr<double> p(new double);  
*(p.get()) = 9.8;  
std::shared_ptr<double> q = p;  
std::cout << *(q.get()) << std::endl;
```

- and the memory is correctly freed once at the end of the program.
- A similar construct is `std::unique_ptr` which only allows a single pointer to a block of memory to exist.
- However, the above only supports `new[]` from C++17 onwards.

Failed allocation of memory

- If you are out of memory, or try to allocate too much memory:

```
size_t bigMemory = 1L << 36;  
double* myBigData = new double[bigMemory];
```

then an exception will be thrown.

- If you don't put any exception-handling code in, then your program will abort, which is probably what you want.
- `size_t` is an unsigned integer big enough to refer to all memory that might be needed on the current architecture.
- So, on a 64-bit system, `size_t` will be 8 bytes (64 bits) in size as the memory addressing uses 64 bits.
- (Note: `1L` is required to make `1L << 36` a 64-bit integer.)

Multi-dimensional arrays (heap)

- What about allocating multi-dimensional arrays on the heap?
- The short answer is that you can't
- The longer answer is that you need to create an `Array` class that deals with the access
- or allocate firstly a contiguous block of memory and then a set of pointers to successive rows of that.
- Either:

```
int *a = new int[N*M];  
int aIJ = a[N*i + j];
```

- Or:

```
int *a = new int[N*M];  
int **b = new int *[M];  
b[0] = a; b[1] = a + N;  
int aIJ = b[i][j];
```

sizeof

- If you need to determine how much memory is needed for an array, you can use `sizeof`:

```
size_t numBytesReqd = sizeof(double) * 100;
```

- This can also be applied to variables:

```
int b = 3;  
size_t sizeB = sizeof(b); // e.g. 4 bytes
```

- Warning, the following does not behave as you might expect for an array:

```
double* a = new double[100];  
size_t aSize = sizeof(a); // 8 bytes (on 64-bit machine)
```

returning the size of a pointer variable.

- The `sizeof` operator returns a variable of type `size_t`, which is used whenever the size of an object in memory is required.

Outline

19 Pointers

20 Memory

21 **Function wrappers**

Function wrappers

- We may have the user choose an ODE solver function at run-time
- Assume we have two or more functions with a consistent interface:

```
double euler(double x, double t, double dt) {  
    return x + dt * y(x,t);  
}  
double rk2(double x, double t, double dt) {  
    // RK2 code here  
    return x + dt * dy;  
}
```

- and we want to call whichever one of these the user has chosen.

Function wrappers ctd

- We could have a switch statement every time we need to call one or the other:

```
double nextX;  
switch(odeSolver) {  
    case ODEsolver::EULER:  
        nextX = euler(x,t,dt);  
        break;  
    case ODEsolver::RK2:  
        nextX = rk2(x,t,dt);  
        break;  
}
```

- However, if this needs to occur multiple times, then we end up with multiple copies of this code, and what happens if we want to add RK4?

Function wrappers

- We would like to store a reference to the function:

```
#include <functional>
std::function<double(double, double, double)> odeSolverFn;
switch(odeSolver) {
    case ODEsolver::EULER:
        odeSolverFn = &euler;
        break;
    case ODEsolver::RK2:
        odeSolverFn = rk2; // & not necessary
        break;
}
```

- This is used as:

```
nextX = odeSolverFn(x,t,dt);
```

Function wrapper cost

- Since the compiler no longer knows which function will be called at compile-time, this introduces an extra redirection call at machine-code level
- This has a very small computational cost.
- Unless you profile your real-world code and find it's expensive, don't worry about the extra cost.
- Note that function wrappers have replaced function *pointers* which were harder to get right, and less powerful.
- In older code, you may see the more confusing:

```
double (*odeSolverFn)(double, double, double) = &euler;
```